



Vidya Vardhini's College of Engineering and Technology
Department of Artificial Intelligence & Data Science

AY: 2024 - 25

Class:	BE	Semester:	VII
Course Code:		Course Name:	BDA

Name of Student:	BARI ANIKIT VINOD
Roll No. :	61
Assignment No.:	3
Title of Assignment:	
Date of Submission:	
Date of Correction:	

Evaluation

Performance Indicator	Max. Marks	Marks Obtained
Completeness	5	4
Demonstrated Knowledge Legibility	3	3
Legibility	2	2
Total	10	9

Performance Indicator	Exceed Expectations (EE)	Meet Expectations (ME)	Below Expectations (BE)
Completeness	5	3-4	1-2
Demonstrated Knowledge Legibility	3	2	1
Legibility	2	1	0

Checked by

Name of Faculty : Ms. Sweety Patil
 Signature :
 Date : 4/9/23

Assignment No. - 3

1) Apply the concept of column family store and graph store of NOSQL architecture patterns on student management systems.

① Column-family store - (modeling principles) -

- design for queries - pick partition keys + clustering keys so the most frequent queries are served by single partition scans.
- denormalize - duplicate data as needed to avoid joins.
- wide rows - store sequential / time + series events as clustering.

Eg.

CREATE TABLE student_profiles (

student_id uuid,

name text,

dob date,

email text,

phone text,

address text,

major text,

PRIMARY KEY (student_id));

② Typical Graph store - (Graph model elements) -

- Nodes : Student, Course, Instructor, Department, Assignment, Group.

- Relationships : :ENROLLED-IN, :ADVISED-BY, :TA-FOR, :PREREQ-OF, :WORKED-WITH, :SUBMITTED.

- Properties : node / edge properties like grade, semester, role, timestamp

Eg. Cypher / Neo4j constructions -

CREATE (s1: Student {id: '123', 'name': 'Alice'}),

(c1: Course {id: 'CS101', 'title': 'Intro to CS'}),

(c2: Course {id: 'CS201', 'title': 'Data Structures'}),

(i1: Instructor {id: 'IT01', 'name': 'Dr. Pav'}));

CREATE (s1) - [e: ENROLLED-IN {semester: '2025-Fall', grade: 'A'}] -> (c1);

CREATE (c1) - [e: PREREQ-OF] -> (c2);

CREATE (s1) - [e: ADVISED-BY] -> (i1);

Q. 2)

Consider the case of Netflix adopting Apache Cassandra (NoSQL). Show how business problems like high availability, real-time, personalization, and global scalability were solved faster, cheaper, and more effectively. State the business drivers and the findings in the implementations.

→ (i) Business drivers for netflix -

- High availability - (Always on service) - Streaming must never go down - down time means instant user churn. User expect 24x7 uptime across all regions, even during Outages.

(ii) Real-time personalization -

- Netflix's value is not just movies but personalized recommendations.
- Need sub-second updates of user activity to adapt recommendations.

(iii) Global scalability -

- Tens or thousands of microservices handle petabytes of data.
- User world wide needs low-latency access data must be across regions.

(iv) Cost-effectiveness -

- Traditional RDBMS sharding and replication across constraints would be too costly and complex.

- Need elastic, commodity-hardware - friendly system.

• Example use cases of cassandra in netflix -

- viewing history - stores every video watched by users.

- user preferences - stores likes/dislike, device activity, profiles.

- Recommendations & search - cassandra + lucene search + ML models

for personalization results.

- billing & accounting events - distributed, durable writes across regions

- fast development cycle - cassandra's schema flexibility reduced engineering overhead.

- cheaper operations - commodity cloud infrastructure + linear scale.

- resilience - survive large-scale AWS outages without downtime.

- community contribution - Netflix created tools like psion.